

Week 2 Exercises

Kaylynn Hiller

2026-01-20

Week 2

```
-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
v dplyr     1.1.4     v readr      2.1.5
v forcats   1.0.1     v stringr    1.6.0
v ggplot2   4.0.0     v tibble     3.3.1
v lubridate 1.9.4     v tidyr     1.3.2
v purrr     1.2.1

-- Conflicts -----
x dplyr::filter() masks stats::filter()
x dplyr::lag()    masks stats::lag()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to beco
```

Exercise 3.2.5

1. In a single pipeline for each condition, find all flights that meet the condition:
 - Had an arrival delay of two or more hours

```
# A tibble: 10,200 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>    <int>        <int>    <dbl>    <int>        <int>
1 2013     1     1      811          630      101       1047        830
2 2013     1     1      848          1835      853      1001       1950
3 2013     1     1      957          733      144       1056        853
4 2013     1     1     1114          900      134       1447       1222
5 2013     1     1     1505         1310      115       1638       1431
6 2013     1     1     1525         1340      105       1831       1626
7 2013     1     1     1549         1445       64       1912       1656
8 2013     1     1     1558         1359      119       1718       1515
```

```

9 2013    1    1    1732      1630      62    2028      1825
10 2013    1    1    1803     1620      103   2008      1750
# i 10,190 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
# tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
# hour <dbl>, minute <dbl>, time_hour <dttm>

```

- Flew to Houston (IAH or HOU)

```

# A tibble: 9,313 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>     <int>       <int>     <dbl>     <int>       <int>
1 2013    1    1      517        515        2     830        819
2 2013    1    1      533        529        4     850        830
3 2013    1    1      623        627       -4     933        932
4 2013    1    1      728        732       -4    1041       1038
5 2013    1    1      739        739        0    1104       1038
6 2013    1    1      908        908        0    1228       1219
7 2013    1    1     1028       1026        2    1350       1339
8 2013    1    1     1044       1045       -1    1352       1351
9 2013    1    1     1114       900       134    1447       1222
10 2013   1    1     1205      1200        5    1503       1505
# i 9,303 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
# tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
# hour <dbl>, minute <dbl>, time_hour <dttm>

```

- Were operated by United, American, or Delta

```

# A tibble: 139,504 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>     <int>       <int>     <dbl>     <int>       <int>
1 2013    1    1      517        515        2     830        819
2 2013    1    1      533        529        4     850        830
3 2013    1    1      542        540        2     923        850
4 2013    1    1      554        600       -6     812        837
5 2013    1    1      554        558       -4     740        728
6 2013    1    1      558        600       -2     753        745
7 2013    1    1      558        600       -2     924        917
8 2013    1    1      558        600       -2     923        937
9 2013    1    1      559        600       -1     941        910
10 2013   1    1      559        600       -1     854        902

```

```
# i 139,494 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dttm>
```

- Departed in summer (July, August, and September)

```
# A tibble: 86,326 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>     <int>        <int>    <dbl>    <int>        <int>
1 2013     7     1       1      2029      212      236        2359
2 2013     7     1       2      2359       3      344        344
3 2013     7     1      29      2245      104      151         1
4 2013     7     1      43      2130      193      322        14
5 2013     7     1      44      2150      174      300        100
6 2013     7     1      46      2051      235      304        2358
7 2013     7     1      48      2001      287      308        2305
8 2013     7     1      58      2155      183      335        43
9 2013     7     1     100      2146      194      327        30
10 2013    7     1     100      2245      135      337        135
# i 86,316 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dttm>
```

- Arrived more than two hours late but didn't leave late

```
# A tibble: 29 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>     <int>        <int>    <dbl>    <int>        <int>
1 2013     1    27     1419      1420      -1      1754        1550
2 2013    10     7     1350      1350       0      1736        1526
3 2013    10     7     1357      1359      -2      1858        1654
4 2013    10    16     657       700      -3      1258        1056
5 2013    11     1     658       700      -2      1329        1015
6 2013     3    18     1844      1847      -3       39        2219
7 2013     4    17     1635      1640      -5      2049        1845
8 2013     4    18     558       600      -2      1149        850
9 2013     4    18     655       700      -5      1213        950
10 2013    5    22     1827      1830      -3      2217        2010
# i 19 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
```

```

# tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
# hour <dbl>, minute <dbl>, time_hour <dttm>

• Were delayed by at least an hour, but made up over 30 minutes in flight

# A tibble: 1,844 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>    <int>        <int>     <dbl>    <int>        <int>
1 2013     1     1      2205          1720      285       46        2040
2 2013     1     1      2326          2130      116      131        18
3 2013     1     3      1503          1221      162      1803      1555
4 2013     1     3      1839          1700       99      2056      1950
5 2013     1     3      1850          1745       65      2148      2120
6 2013     1     3      1941          1759      102      2246      2139
7 2013     1     3      1950          1845       65      2228      2227
8 2013     1     3      2015          1915       60      2135      2111
9 2013     1     3      2257          2000      177       45      2224
10 2013    1     4      1917          1700      137      2135      1950
# i 1,834 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
# tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
# hour <dbl>, minute <dbl>, time_hour <dttm>

```

2. Sort flights to find the flights with the longest departure delays. Find the flights that left earliest in the morning.

```

# A tibble: 336,776 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>    <int>        <int>     <dbl>    <int>        <int>
1 2013     1     9      641          900      1301      1242      1530
2 2013     6    15     1432         1935      1137      1607      2120
3 2013     1    10     1121         1635      1126      1239      1810
4 2013     9    20     1139         1845      1014      1457      2210
5 2013     7    22      845         1600      1005      1044      1815
6 2013     4    10     1100         1900      960       1342      2211
7 2013     3    17     2321        810       911       135       1020
8 2013     6    27      959         1900      899       1236      2226
9 2013     7    22     2257        759       898       121       1026
10 2013    12     5      756        1700      896       1058      2020
# i 336,766 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
# tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
# hour <dbl>, minute <dbl>, time_hour <dttm>

```

```

# A tibble: 336,776 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>     <int>          <int>     <dbl>     <int>          <int>
1 2013     1    13      1        2249       72     108        2357
2 2013     1    31      1        2100      181     124        2225
3 2013    11    13      1        2359       2     442        440
4 2013    12    16      1        2359       2     447        437
5 2013    12    20      1        2359       2     430        440
6 2013    12    26      1        2359       2     437        440
7 2013    12    30      1        2359       2     441        437
8 2013     2    11      1        2100      181     111        2225
9 2013     2    24      1        2245       76     121        2354
10 2013     3     8      1        2355       6     431        440
# i 336,766 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dttm>

```

3. Sort flights to find the fastest flights. (Hint: Try including a math calculation inside of your function.)

```

# A tibble: 336,776 x 20
  speed year month   day dep_time sched_dep_time dep_delay arr_time
  <dbl> <int> <int> <int>          <int>     <dbl>     <int>
1 1132. 2013     5    25      1709       1700       9     1923
2 1047. 2013     7     2      1558       1513      45     1745
3 1043. 2013     5    13      2040       2025      15     2225
4 1032. 2013     3    23      1914       1910       4     2045
5  952. 2013     1    12      1559       1600      -1     1849
6  908. 2013    11    17      650        655      -5     1059
7  897. 2013     2    21      2355       2358      -3      412
8  896. 2013    11    17      759        800      -1     1212
9  892. 2013    11    16      2003       1925      38      17
10 892. 2013    11    16      2349       2359     -10      402
# i 336,766 more rows
# i 12 more variables: sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
#   flight <int>, tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>,
#   distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dttm>

```

4. Was there a flight on every day of 2013? Yes

```
# A tibble: 1 x 1
```

```

num_days
<int>
1      365

```

5. Which flights traveled the farthest distance? Which traveled the least distance?

```

# A tibble: 336,776 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>    <int>          <int>     <dbl>    <int>          <int>
1 2013     1     1      857           900      -3     1516          1530
2 2013     1     2      909           900       9     1525          1530
3 2013     1     3      914           900      14     1504          1530
4 2013     1     4      900           900       0     1516          1530
5 2013     1     5      858           900      -2     1519          1530
6 2013     1     6     1019           900      79     1558          1530
7 2013     1     7     1042           900     102     1620          1530
8 2013     1     8      901           900       1     1504          1530
9 2013     1     9      641           900     1301     1242          1530
10 2013    1    10      859           900      -1     1449          1530
# i 336,766 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dttm>

# A tibble: 336,776 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>    <int>          <int>     <dbl>    <int>          <int>
1 2013     7    27      NA           106      NA      NA          245
2 2013     1     3     2127          2129      -2     2222          2224
3 2013     1     4     1240          1200      40     1333          1306
4 2013     1     4     1829          1615     134     1937          1721
5 2013     1     4     2128          2129      -1     2218          2224
6 2013     1     5     1155          1200      -5     1241          1306
7 2013     1     6     2125          2129      -4     2224          2224
8 2013     1     7     2124          2129      -5     2212          2224
9 2013     1     8     2127          2130      -3     2304          2225
10 2013    1     9     2126          2129      -3     2217          2224
# i 336,766 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dttm>

```

6. Does it matter what order you used filter() and arrange() if you're using both? Why/why not? Think about the results and how much work the functions would have to do.

It does not matter. Filter pulls out rows that meet a condition but does not change the ordering. Arrange changes the ordering but leaves all rows in. Since they perform separate actions, the order does not matter.

Exercise 3.3.5

1. Compare dep_time, sched_dep_time, and dep_delay. How would you expect those three numbers to be related?

I would expect the scheduled departure time (sched_dep_time) plus the delay (dep_delay) to equal the actual departure time (dep_time).

```
# A tibble: 336,776 x 3
  dep_time sched_dep_time dep_delay
  <int>       <int>     <dbl>
1      517         515      2
2      533         529      4
3      542         540      2
4      544         545     -1
5      554         600     -6
6      554         558     -4
7      555         600     -5
8      557         600     -3
9      557         600     -3
10     558         600     -2
# i 336,766 more rows
```

2. Brainstorm as many ways as possible to select dep_time, dep_delay, arr_time, and arr_delay from flights.

```
# A tibble: 336,776 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>       <int>     <int>     <dbl>     <int>       <int>
1 2013     1     1      517         515      2        830       819
2 2013     1     1      533         529      4        850       830
3 2013     1     1      542         540      2        923       850
4 2013     1     1      544         545     -1       1004      1022
5 2013     1     1      554         600     -6       812       837
6 2013     1     1      554         558     -4       740       728
7 2013     1     1      555         600     -5       913       854
```

```

8 2013 1 1 557 600 -3 709 723
9 2013 1 1 557 600 -3 838 846
10 2013 1 1 558 600 -2 753 745
# i 336,766 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
# tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
# hour <dbl>, minute <dbl>, time_hour <dttm>

# A tibble: 336,776 x 4
  dep_time dep_delay arr_time arr_delay
  <int>     <dbl>    <int>     <dbl>
1      517        2     830       11
2      533        4     850       20
3      542        2     923       33
4      544       -1    1004      -18
5      554       -6     812      -25
6      554       -4     740       12
7      555       -5     913       19
8      557       -3     709      -14
9      557       -3     838       -8
10     558       -2     753        8
# i 336,766 more rows

# A tibble: 336,776 x 6
  dep_time sched_dep_time dep_delay arr_time sched_arr_time arr_delay
  <int>          <int>     <dbl>    <int>          <int>     <dbl>
1      517          515        2     830          819       11
2      533          529        4     850          830       20
3      542          540        2     923          850       33
4      544          545       -1    1004         1022      -18
5      554          600       -6     812          837      -25
6      554          558       -4     740          728       12
7      555          600       -5     913          854       19
8      557          600       -3     709          723      -14
9      557          600       -3     838          846       -8
10     558          600       -2     753          745        8
# i 336,766 more rows

# A tibble: 336,776 x 4
  dep_time dep_delay arr_time arr_delay
  <int>     <dbl>    <int>     <dbl>
1      517        2     830       11

```

```

2      533        4     850       20
3      542        2     923       33
4      544       -1    1004      -18
5      554       -6     812      -25
6      554       -4     740       12
7      555       -5     913       19
8      557       -3     709      -14
9      557       -3     838       -8
10     558       -2     753        8
# i 336,766 more rows

# A tibble: 336,776 x 4
  dep_time dep_delay arr_time arr_delay
  <int>     <dbl>    <int>     <dbl>
1      517         2     830        11
2      533         4     850       20
3      542         2     923       33
4      544        -1    1004      -18
5      554        -6     812      -25
6      554        -4     740       12
7      555        -5     913       19
8      557        -3     709      -14
9      557        -3     838       -8
10     558        -2     753        8
# i 336,766 more rows

```

3. What happens if you specify the name of the same variable multiple times in a select() call?

```

# A tibble: 336,776 x 1
  year
  <int>
1 2013
2 2013
3 2013
4 2013
5 2013
6 2013
7 2013
8 2013
9 2013
10 2013
# i 336,766 more rows

```

You just get the variable once.

4. What does the any_of() function do? Why might it be helpful in conjunction with this vector?

```
# A tibble: 336,776 x 5
  year month   day dep_delay arr_delay
  <int> <int> <int>     <dbl>     <dbl>
1 2013     1     1       2      11
2 2013     1     1       4      20
3 2013     1     1       2      33
4 2013     1     1      -1     -18
5 2013     1     1      -6     -25
6 2013     1     1      -4      12
7 2013     1     1      -5      19
8 2013     1     1      -3     -14
9 2013     1     1      -3      -8
10 2013    1     1      -2       8
# i 336,766 more rows
```

The any_of() function in conjunction with the select function selects any variables listed. If you set specific variables to select outside the pipeline, the code is cleaner and easier to change.

5. Does the result of running the following code surprise you? How do the select helpers deal with upper and lower case by default? How can you change that default?

```
# A tibble: 336,776 x 6
  dep_time sched_dep_time arr_time sched_arr_time air_time time_hour
  <int>        <int>     <int>        <int>     <dbl> <dttm>
1     517          515     830         819      227 2013-01-01 05:00:00
2     533          529     850         830      227 2013-01-01 05:00:00
3     542          540     923         850      160 2013-01-01 05:00:00
4     544          545    1004        1022      183 2013-01-01 05:00:00
5     554          600     812         837      116 2013-01-01 06:00:00
6     554          558     740         728      150 2013-01-01 05:00:00
7     555          600     913         854      158 2013-01-01 06:00:00
8     557          600     709         723       53 2013-01-01 06:00:00
9     557          600     838         846      140 2013-01-01 06:00:00
10    558          600     753         745      138 2013-01-01 06:00:00
# i 336,766 more rows
```

The result did not surprise me. By default, select helpers treat upper and lower case the same. You can change this by specifying with ignore.case = FALSE.

6. Rename air_time to air_time_min to indicate units of measurement and move it to the beginning of the data frame.

```
# A tibble: 336,776 x 19
  air_time_min year month   day dep_time sched_dep_time dep_delay arr_time
  <dbl> <int> <int> <int> <int> <int> <dbl> <int>
1       227  2013     1     1    517      515      2     830
2       227  2013     1     1    533      529      4     850
3       160  2013     1     1    542      540      2     923
4       183  2013     1     1    544      545     -1    1004
5       116  2013     1     1    554      600     -6     812
6       150  2013     1     1    554      558     -4     740
7       158  2013     1     1    555      600     -5     913
8        53  2013     1     1    557      600     -3     709
9       140  2013     1     1    557      600     -3     838
10      138  2013     1     1    558      600     -2     753
# i 336,766 more rows
# i 11 more variables: sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
#   flight <int>, tailnum <chr>, origin <chr>, dest <chr>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dttm>
```

7. Why doesn't the following work, and what does the error mean?

We did not select the arr_delay variable in the select function so it is not available to arrange. Select keeps only the columns listed.

Exercise 3.5.7

- Which carrier has the worst average delays? Challenge: can you disentangle the effects of bad airports vs. bad carriers? Why/why not? (Hint: think about...)

```
# A tibble: 16 x 2
  carrier avg_dep_del
  <chr>      <dbl>
1 F9          20.2
2 EV          20.0
3 YV          19.0
4 FL          18.7
5 WN          17.7
6 9E          16.7
7 B6          13.0
8 VX          12.9
9 OO          12.6
```

```
10 UA      12.1
11 MQ      10.6
12 DL      9.26
13 AA      8.59
14 AS      5.80
15 HA      4.90
16 US      3.78
```

```
# A tibble: 1 x 4
  carrier dest  origin     n
  <chr>   <chr> <chr> <int>
1 F9      DEN    LGA     685
```

```
# A tibble: 5 x 3
  carrier dest     n
  <chr>   <chr> <int>
1 B6      DEN     338
2 DL      DEN    1043
3 F9      DEN     685
4 UA      DEN    3796
5 WN      DEN    1404
```

`summarise()` has grouped output by 'carrier', 'dest'. You can override using the ` `.groups` argument.

```
# A tibble: 8 x 5
# Groups:   carrier, dest [5]
  carrier dest  origin     n avg_dd
  <chr>   <chr> <chr> <int>  <dbl>
1 DL      DEN    LGA     678   9.82
2 UA      DEN    LGA    1626  10.6
3 UA      DEN    EWR    2170  14.1
4 WN      DEN    LGA     715  15.4
5 DL      DEN    JFK     365  16.6
6 F9      DEN    LGA     685  20.2
7 B6      DEN    JFK     338  23.9
8 WN      DEN    EWR     689  24.4
```

The carrier with the worst average delays is Frontier airlines. However, they are headquartered in Denver and this data set only includes flights from LGA to DEN. This means there's

no variance in F9 across airports since there is only one flight path being recorded. We cannot disentangle the effects of bad airports from this data set and would have to make a lot of assumptions in order to do so.

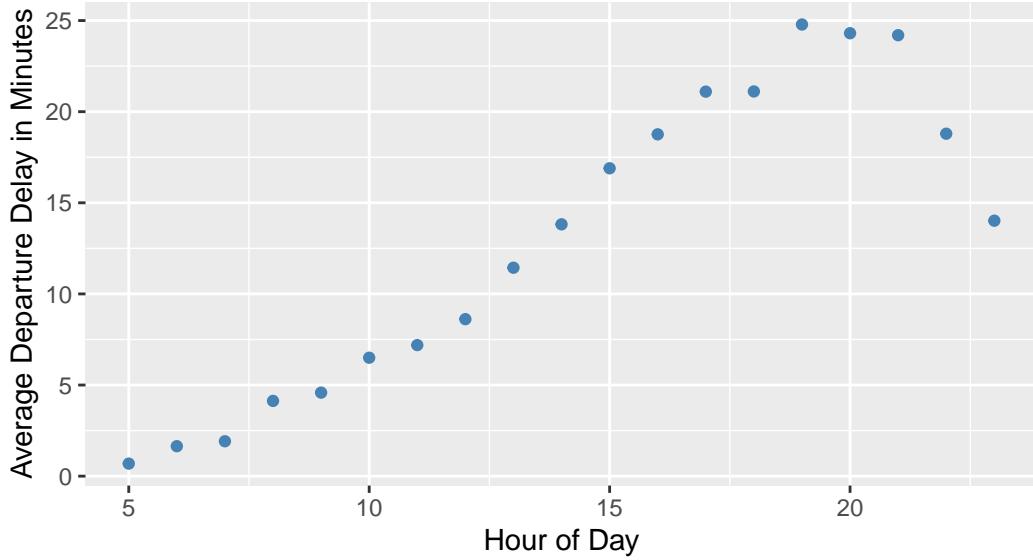
2. Find the flights that are most delayed upon departure to each destination.

```
# A tibble: 105 x 19
# Groups:   dest [105]
  dest carrier dep_delay year month   day dep_time sched_dep_time arr_time
  <chr> <chr>     <dbl> <int> <int> <int>     <int>           <int>     <int>
1 HNL  HA        1301  2013     1     9      641            900     1242
2 CMH  MQ        1137  2013     6    15     1432           1935     1607
3 ORD  MQ        1126  2013     1    10     1121           1635     1239
4 SFO  AA        1014  2013     9    20     1139           1845     1457
5 CVG  MQ        1005  2013     7    22      845            1600     1044
6 TPA  DL        960   2013     4    10     1100           1900     1342
7 MSP  DL        911   2013     3    17     2321            810      135
8 PDX  DL        899   2013     6    27      959           1900     1236
9 ATL  DL        898   2013     7    22     2257            759      121
10 MIA  AA       896   2013    12     5      756           1700     1058
# i 95 more rows
# i 10 more variables: sched_arr_time <int>, arr_delay <dbl>, flight <int>,
#   tailnum <chr>, origin <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
#   minute <dbl>, time_hour <dttm>
```

3. How do delays vary over the course of the day? Illustrate your answer with a plot.

Average Departure Delays Over a 24 Hour Period

Using US Bureau of Transportation Statistics Data on Flights Departing from



Delays do vary.

4. What happens if you supply a negative n to slice_min() and friends?

```
# A tibble: 336,776 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>    <int>        <int>     <dbl>    <int>        <int>
1 2013     12     7    2040        2123      -43        40        2352
2 2013      2     3    2022        2055      -33       2240        2338
3 2013     11    10    1408        1440      -32       1549       1559
4 2013      1    11    1900        1930      -30       2233       2243
5 2013      1    29    1703        1730      -27       1947       1957
6 2013      8     9     729         755      -26       1002       955
7 2013     10    23    1907        1932      -25       2143       2143
8 2013      3    30    2030        2055      -25       2213       2250
9 2013      3     2    1431        1455      -24       1601       1631
10 2013      5     5    934         958      -24       1225       1309
# i 336,766 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dttm>

# A tibble: 336,776 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
```

```

<int> <int> <int>      <int>      <int>      <dbl>      <int>      <int>
1 2013     12     7     2040     2123     -43       40     2352
2 2013     2     3     2022     2055     -33      2240     2338
3 2013     11    10     1408     1440     -32      1549     1559
4 2013     1    11     1900     1930     -30      2233     2243
5 2013     1    29     1703     1730     -27      1947     1957
6 2013     8     9     729      755     -26      1002      955
7 2013    10    23     1907     1932     -25      2143     2143
8 2013     3    30     2030     2055     -25      2213     2250
9 2013     3     2     1431     1455     -24      1601     1631
10 2013    5     5     934      958     -24     1225     1309
# i 336,766 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dttm>

```

The result is listing the rows based on the smallest to largest dep_delay value. This result is the same as arranging the values by dep_delay.

5. Explain what count() does in terms of the dplyr verbs you just learned. What does the sort argument to count() do?

```
# A tibble: 3 x 2
  origin     n
  <chr>   <int>
1 EWR     120835
2 JFK     111279
3 LGA     104662
```

```
# A tibble: 3 x 2
  origin     n
  <chr>   <int>
1 EWR     120835
2 JFK     111279
3 LGA     104662
```

Count allows you to count the unique values of one or more variables similar to group_by and summarize. The sort options causes the results to be presented in descending order.

6. Suppose we have the following tiny data frame:

- a. Write down what you think the output will look like, then check if you were correct, and describe what group_by() does.

```
# A tibble: 5 x 3
# Groups:   y [2]
  x     y     z
<int> <chr> <chr>
1     1     a     K
2     2     b     K
3     3     a     L
4     4     a     L
5     5     b     K
```

This command should group by the two groups in y.

- b. Write down what you think the output will look like, then check if you were correct, and describe what arrange() does. Also, comment on how it's different from the group_by() in part (a).

```
# A tibble: 5 x 3
  x     y     z
<int> <chr> <chr>
1     1     a     K
2     3     a     L
3     4     a     L
4     2     b     K
5     5     b     K
```

This command should arrange the results by the y column alphabetically. This means it should be list the results as a,a,a,b,b. The row order changed in using arrange() and not group_by().

- c. Write down what you think the output will look like, then check if you were correct, and describe what the pipeline does.

```
# A tibble: 2 x 2
  y      mean_x
<chr>    <dbl>
1 a        2.67
2 b        3.5
```

The result should be a summary of the mean of x based on the two groups of y. The pipeline passes the group_by() function along into the summarize() function so the means are groups by y values.

- d. Write down what you think the output will look like, then check if you were correct, and describe what the pipeline does. Then, comment on what the message says.

`summarise()` has grouped output by 'y'. You can override using the `groups` argument.

```
# A tibble: 3 x 3
# Groups:   y [2]
  y     z     mean_x
  <chr> <chr>   <dbl>
1 a     K         1
2 a     L         3.5
3 b     K         3.5
```

The result should be a summary of the mean of x based on the three groups of y and z (a K, a L, a K). This pipeline passes the group_by() function along into the summarize() function so the means are groups by y and z values. The message means that the results are grouped by y values first and the grouping can be overridden using the .groups argument.

- e. Write down what you think the output will look like, then check if you were correct, and describe what the pipeline does. How is the output different from the one in part (d)?

```
# A tibble: 3 x 3
  y     z     mean_x
  <chr> <chr>   <dbl>
1 a     K         1
2 a     L         3.5
3 b     K         3.5
```

The answer should be the same except the results are not grouped.

- f. Write down what you think the outputs will look like, then check if you were correct, and describe what each pipeline does. How are the outputs of the two pipelines different?

`summarise()` has grouped output by 'y'. You can override using the `groups` argument.

```
# A tibble: 3 x 3
# Groups:   y [2]
  y     z     mean_x
  <chr> <chr>   <dbl>
```

```

1 a      K      1
2 a      L      3.5
3 b      K      3.5

# A tibble: 5 x 4
# Groups:   y, z [3]
  x     y     z   mean_x
<int> <chr> <chr>    <dbl>
1     1 a     K      1
2     2 b     K      3.5
3     3 a     L      3.5
4     4 a     L      3.5
5     5 b     K      3.5

```

The first pipeline calculated the means for the three combinations of the y and z column values. The second pipeline creates a new column where these means are stored. The output of the first pipeline is three rows and five rows for the second.