

# Week 2 Exercises

Kaylynn Hiller

2026-01-16

## Week 2

```
-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
v dplyr     1.1.4     v readr     2.1.5
v forcats   1.0.1     v stringr   1.6.0
v ggplot2   4.0.0     v tibble    3.3.1
v lubridate 1.9.4     v tidyr    1.3.2
v purrr    1.2.1

-- Conflicts -----
x dplyr::filter() masks stats::filter()
x dplyr::lag()    masks stats::lag()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to beco
```

### Exercise 3.2.5

1. In a single pipeline for each condition, find all flights that meet the condition:
  - Had an arrival delay of two or more hours

```
# A tibble: 10,200 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>    <int>        <int>    <dbl>    <int>        <int>
1 2013     1     1      811          630      101       1047        830
2 2013     1     1      848          1835      853      1001       1950
3 2013     1     1      957          733      144       1056        853
4 2013     1     1     1114          900      134       1447       1222
5 2013     1     1     1505         1310      115       1638       1431
6 2013     1     1     1525         1340      105       1831       1626
7 2013     1     1     1549         1445       64       1912       1656
8 2013     1     1     1558         1359      119       1718       1515
```

```

9 2013    1    1    1732      1630      62    2028      1825
10 2013    1    1    1803     1620      103   2008      1750
# i 10,190 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
# tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
# hour <dbl>, minute <dbl>, time_hour <dttm>

```

- Flew to Houston (IAH or HOU)

```

# A tibble: 9,313 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>     <int>       <int>     <dbl>     <int>       <int>
1 2013    1    1      517        515        2     830        819
2 2013    1    1      533        529        4     850        830
3 2013    1    1      623        627       -4     933        932
4 2013    1    1      728        732       -4    1041       1038
5 2013    1    1      739        739        0    1104       1038
6 2013    1    1      908        908        0    1228       1219
7 2013    1    1     1028       1026        2    1350       1339
8 2013    1    1     1044       1045       -1    1352       1351
9 2013    1    1     1114       900       134    1447       1222
10 2013   1    1     1205      1200        5    1503       1505
# i 9,303 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
# tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
# hour <dbl>, minute <dbl>, time_hour <dttm>

```

- Were operated by United, American, or Delta

```

# A tibble: 139,504 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>     <int>       <int>     <dbl>     <int>       <int>
1 2013    1    1      517        515        2     830        819
2 2013    1    1      533        529        4     850        830
3 2013    1    1      542        540        2     923        850
4 2013    1    1      554        600       -6     812        837
5 2013    1    1      554        558       -4     740        728
6 2013    1    1      558        600       -2     753        745
7 2013    1    1      558        600       -2     924        917
8 2013    1    1      558        600       -2     923        937
9 2013    1    1      559        600       -1     941        910
10 2013   1    1      559        600       -1     854        902

```

```
# i 139,494 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dttm>
```

- Departed in summer (July, August, and September)

```
# A tibble: 86,326 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>     <int>        <int>    <dbl>    <int>        <int>
1 2013     7     1       1      2029      212      236        2359
2 2013     7     1       2      2359       3      344        344
3 2013     7     1      29      2245      104      151         1
4 2013     7     1      43      2130      193      322        14
5 2013     7     1      44      2150      174      300        100
6 2013     7     1      46      2051      235      304        2358
7 2013     7     1      48      2001      287      308        2305
8 2013     7     1      58      2155      183      335        43
9 2013     7     1     100      2146      194      327        30
10 2013    7     1     100      2245      135      337        135
# i 86,316 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dttm>
```

- Arrived more than two hours late but didn't leave late

```
# A tibble: 29 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>     <int>        <int>    <dbl>    <int>        <int>
1 2013     1    27     1419      1420      -1      1754        1550
2 2013    10     7     1350      1350       0      1736        1526
3 2013    10     7     1357      1359      -2      1858        1654
4 2013    10    16     657       700      -3      1258        1056
5 2013    11     1     658       700      -2      1329        1015
6 2013     3    18     1844      1847      -3       39        2219
7 2013     4    17     1635      1640      -5      2049        1845
8 2013     4    18     558       600      -2      1149        850
9 2013     4    18     655       700      -5      1213        950
10 2013    5    22     1827      1830      -3      2217        2010
# i 19 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
```

```

# tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
# hour <dbl>, minute <dbl>, time_hour <dttm>

• Were delayed by at least an hour, but made up over 30 minutes in flight

# A tibble: 1,844 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>    <int>        <int>     <dbl>    <int>        <int>
1 2013     1     1      2205          1720      285       46        2040
2 2013     1     1      2326          2130      116      131        18
3 2013     1     3      1503          1221      162      1803      1555
4 2013     1     3      1839          1700       99      2056      1950
5 2013     1     3      1850          1745       65      2148      2120
6 2013     1     3      1941          1759      102      2246      2139
7 2013     1     3      1950          1845       65      2228      2227
8 2013     1     3      2015          1915       60      2135      2111
9 2013     1     3      2257          2000      177       45      2224
10 2013    1     4      1917          1700      137      2135      1950
# i 1,834 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
# tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
# hour <dbl>, minute <dbl>, time_hour <dttm>

```

2. Sort flights to find the flights with the longest departure delays. Find the flights that left earliest in the morning.

```

# A tibble: 336,776 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>    <int>        <int>     <dbl>    <int>        <int>
1 2013     1     9      641          900      1301      1242      1530
2 2013     6    15     1432         1935      1137      1607      2120
3 2013     1    10     1121         1635      1126      1239      1810
4 2013     9    20     1139         1845      1014      1457      2210
5 2013     7    22      845         1600      1005      1044      1815
6 2013     4    10     1100         1900      960       1342      2211
7 2013     3    17     2321        810       911       135       1020
8 2013     6    27      959         1900      899       1236      2226
9 2013     7    22     2257        759       898       121       1026
10 2013    12     5      756        1700      896       1058      2020
# i 336,766 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
# tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
# hour <dbl>, minute <dbl>, time_hour <dttm>

```

```

# A tibble: 336,776 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>     <int>          <int>     <dbl>     <int>          <int>
1 2013     1    13      1        2249       72     108        2357
2 2013     1    31      1        2100      181     124        2225
3 2013    11    13      1        2359       2     442        440
4 2013    12    16      1        2359       2     447        437
5 2013    12    20      1        2359       2     430        440
6 2013    12    26      1        2359       2     437        440
7 2013    12    30      1        2359       2     441        437
8 2013     2    11      1        2100      181     111        2225
9 2013     2    24      1        2245       76     121        2354
10 2013     3     8      1        2355       6     431        440
# i 336,766 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dttm>

```

3. Sort flights to find the fastest flights. (Hint: Try including a math calculation inside of your function.)

```

# A tibble: 336,776 x 20
  speed year month   day dep_time sched_dep_time dep_delay arr_time
  <dbl> <int> <int> <int>          <int>     <dbl>     <int>
1 1132. 2013     5    25      1709       1700       9     1923
2 1047. 2013     7     2      1558       1513      45     1745
3 1043. 2013     5    13      2040       2025      15     2225
4 1032. 2013     3    23      1914       1910       4     2045
5  952. 2013     1    12      1559       1600      -1     1849
6  908. 2013    11    17      650        655      -5     1059
7  897. 2013     2    21      2355       2358      -3      412
8  896. 2013    11    17      759        800      -1     1212
9  892. 2013    11    16      2003       1925      38      17
10 892. 2013    11    16      2349       2359     -10      402
# i 336,766 more rows
# i 12 more variables: sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
#   flight <int>, tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>,
#   distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dttm>

```

4. Was there a flight on every day of 2013? Yes

```
# A tibble: 1 x 1
```

```

num_days
<int>
1      365

```

5. Which flights traveled the farthest distance? Which traveled the least distance?

```

# A tibble: 336,776 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>     <int>        <int>    <dbl>    <int>        <int>
1 2013     1     1      857          900     -3     1516        1530
2 2013     1     2      909          900      9     1525        1530
3 2013     1     3      914          900     14     1504        1530
4 2013     1     4      900          900      0     1516        1530
5 2013     1     5      858          900     -2     1519        1530
6 2013     1     6     1019          900     79     1558        1530
7 2013     1     7     1042          900    102     1620        1530
8 2013     1     8      901          900      1     1504        1530
9 2013     1     9      641          900    1301     1242        1530
10 2013    1    10      859          900     -1     1449        1530
# i 336,766 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dttm>

# A tibble: 336,776 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>     <int>        <int>    <dbl>    <int>        <int>
1 2013     7    27       NA          106      NA      NA        245
2 2013     1     3     2127         2129     -2     2222        2224
3 2013     1     4     1240         1200     40     1333        1306
4 2013     1     4     1829         1615     134     1937        1721
5 2013     1     4     2128         2129     -1     2218        2224
6 2013     1     5     1155         1200     -5     1241        1306
7 2013     1     6     2125         2129     -4     2224        2224
8 2013     1     7     2124         2129     -5     2212        2224
9 2013     1     8     2127         2130     -3     2304        2225
10 2013    1     9     2126         2129     -3     2217        2224
# i 336,766 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dttm>

```

6. Does it matter what order you used filter() and arrange() if you're using both? Why/why not? Think about the results and how much work the functions would have to do.

It does not matter. Filter pulls out rows that meet a condition but does not change the ordering. Arrange changes the ordering but leaves all rows in. Since they perform separate functions, the order does not matter.

### Exercise 3.3.5

1. Compare dep\_time, sched\_dep\_time, and dep\_delay. How would you expect those three numbers to be related?

I would expect the scheduled departure time (sched\_dep\_time) plus the delay (dep\_delay) to equal the actual departure time (dep\_time).

```
# A tibble: 336,776 x 3
  dep_time sched_dep_time dep_delay
  <int>        <int>     <dbl>
1      517          515      2
2      533          529      4
3      542          540      2
4      544          545     -1
5      554          600     -6
6      554          558     -4
7      555          600     -5
8      557          600     -3
9      557          600     -3
10     558          600     -2
# i 336,766 more rows
```

2. Brainstorm as many ways as possible to select dep\_time, dep\_delay, arr\_time, and arr\_delay from flights.

```
# A tibble: 336,776 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>    <int>        <int>     <dbl>    <int>        <int>
1 2013     1     1      517          515      2       830        819
2 2013     1     1      533          529      4       850        830
3 2013     1     1      542          540      2       923        850
4 2013     1     1      544          545     -1      1004       1022
5 2013     1     1      554          600     -6       812        837
6 2013     1     1      554          558     -4       740        728
7 2013     1     1      555          600     -5       913        854
8 2013     1     1      557          600     -3       709        723
```

```

9 2013     1     1      557          600       -3      838        846
10 2013     1     1      558          600       -2      753        745
# i 336,766 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dttm>

# A tibble: 336,776 x 4
  dep_time dep_delay arr_time arr_delay
  <int>     <dbl>    <int>     <dbl>
1     517         2     830       11
2     533         4     850       20
3     542         2     923       33
4     544        -1    1004      -18
5     554        -6     812      -25
6     554        -4     740       12
7     555        -5     913       19
8     557        -3     709      -14
9     557        -3     838       -8
10    558        -2     753        8
# i 336,766 more rows

# A tibble: 336,776 x 6
  dep_time sched_dep_time dep_delay arr_time sched_arr_time arr_delay
  <int>        <int>     <dbl>    <int>        <int>     <dbl>
1     517        515         2     830        819       11
2     533        529         4     850        830       20
3     542        540         2     923        850       33
4     544        545        -1    1004       1022      -18
5     554        600        -6     812        837      -25
6     554        558        -4     740        728       12
7     555        600        -5     913        854       19
8     557        600        -3     709        723      -14
9     557        600        -3     838        846       -8
10    558        600        -2     753        745        8
# i 336,766 more rows

# A tibble: 336,776 x 4
  dep_time dep_delay arr_time arr_delay
  <int>     <dbl>    <int>     <dbl>
1     517         2     830       11
2     533         4     850       20

```

```

3      542        2     923       33
4      544       -1    1004      -18
5      554       -6     812      -25
6      554       -4     740       12
7      555       -5     913       19
8      557       -3     709      -14
9      557       -3     838       -8
10     558       -2     753        8
# i 336,766 more rows

```

```

# A tibble: 336,776 x 4
  dep_time dep_delay arr_time arr_delay
  <int>     <dbl>    <int>     <dbl>
1      517        2     830       11
2      533        4     850       20
3      542        2     923       33
4      544       -1    1004      -18
5      554       -6     812      -25
6      554       -4     740       12
7      555       -5     913       19
8      557       -3     709      -14
9      557       -3     838       -8
10     558       -2     753        8
# i 336,766 more rows

```

3. What happens if you specify the name of the same variable multiple times in a `select()` call?

```

# A tibble: 336,776 x 1
  year
  <int>
1 2013
2 2013
3 2013
4 2013
5 2013
6 2013
7 2013
8 2013
9 2013
10 2013
# i 336,766 more rows

```

You just get the variable once.

4. What does the any\_of() function do? Why might it be helpful in conjunction with this vector?

```
# A tibble: 336,776 x 5
  year month   day dep_delay arr_delay
  <int> <int> <int>     <dbl>     <dbl>
1 2013     1     1       2      11
2 2013     1     1       4      20
3 2013     1     1       2      33
4 2013     1     1      -1     -18
5 2013     1     1      -6     -25
6 2013     1     1      -4      12
7 2013     1     1      -5      19
8 2013     1     1      -3     -14
9 2013     1     1      -3      -8
10 2013    1     1      -2       8
# i 336,766 more rows
```

The any\_of() function in conjunction with the select function selects any variables listed. If you set specific variables to select outside the pipeline, the code is cleaner and easier to change.

5. Does the result of running the following code surprise you? How do the select helpers deal with upper and lower case by default? How can you change that default?

```
# A tibble: 336,776 x 6
  dep_time sched_dep_time arr_time sched_arr_time air_time time_hour
  <int>        <int>     <int>        <int>     <dbl> <dttm>
1     517          515     830         819      227 2013-01-01 05:00:00
2     533          529     850         830      227 2013-01-01 05:00:00
3     542          540     923         850      160 2013-01-01 05:00:00
4     544          545    1004        1022      183 2013-01-01 05:00:00
5     554          600     812         837      116 2013-01-01 06:00:00
6     554          558     740         728      150 2013-01-01 05:00:00
7     555          600     913         854      158 2013-01-01 06:00:00
8     557          600     709         723       53 2013-01-01 06:00:00
9     557          600     838         846      140 2013-01-01 06:00:00
10    558          600     753         745      138 2013-01-01 06:00:00
# i 336,766 more rows
```

The result did not surprise me. By default, select helpers treat upper and lower case the same. You can change this by specifying with ignore.case = FALSE.

6. Rename air\_time to air\_time\_min to indicate units of measurement and move it to the beginning of the data frame.

```
# A tibble: 336,776 x 19
  air_time_min year month   day dep_time sched_dep_time dep_delay arr_time
  <dbl> <int> <int> <int> <int> <int> <dbl> <int>
1     227 2013     1     1    517      515     2    830
2     227 2013     1     1    533      529     4    850
3     160 2013     1     1    542      540     2    923
4     183 2013     1     1    544      545    -1   1004
5     116 2013     1     1    554      600    -6    812
6     150 2013     1     1    554      558    -4    740
7     158 2013     1     1    555      600    -5    913
8      53 2013     1     1    557      600    -3    709
9     140 2013     1     1    557      600    -3    838
10    138 2013     1     1    558      600    -2    753
# i 336,766 more rows
# i 11 more variables: sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
#   flight <int>, tailnum <chr>, origin <chr>, dest <chr>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dttm>
```

7. Why doesn't the following work, and what does the error mean?

We did not select that variable in the select function so it is not available to arrange.

### Exercise 3.5.7

- Which carrier has the worst average delays? Challenge: can you disentangle the effects of bad airports vs. bad carriers? Why/why not? (Hint: think about...)

```
# A tibble: 16 x 2
  carrier avg_dep_del
  <chr>        <dbl>
1 F9          20.2
2 EV          20.0
3 YV          19.0
4 FL          18.7
5 WN          17.7
6 9E          16.7
7 B6          13.0
8 VX          12.9
9 OO          12.6
10 UA         12.1
```

```

11 MQ          10.6
12 DL          9.26
13 AA          8.59
14 AS          5.80
15 HA          4.90
16 US          3.78

# A tibble: 1 x 4
  carrier dest   origin     n
  <chr>    <chr> <chr> <int>
1 F9        DEN    LGA      685

# A tibble: 5 x 3
  carrier dest     n
  <chr>    <chr> <int>
1 B6       DEN     338
2 DL       DEN    1043
3 F9       DEN     685
4 UA       DEN    3796
5 WN       DEN    1404

`summarise()` has grouped output by 'carrier', 'dest'. You can override using
the `.`groups` argument.

# A tibble: 8 x 5
# Groups:   carrier, dest [5]
  carrier dest   origin     n avg_dd
  <chr>    <chr> <chr> <int>  <dbl>
1 DL       DEN    LGA      678   9.82
2 UA       DEN    LGA     1626  10.6
3 UA       DEN    EWR     2170  14.1
4 WN       DEN    LGA      715  15.4
5 DL       DEN    JFK      365  16.6
6 F9       DEN    LGA      685  20.2
7 B6       DEN    JFK      338  23.9
8 WN       DEN    EWR      689  24.4

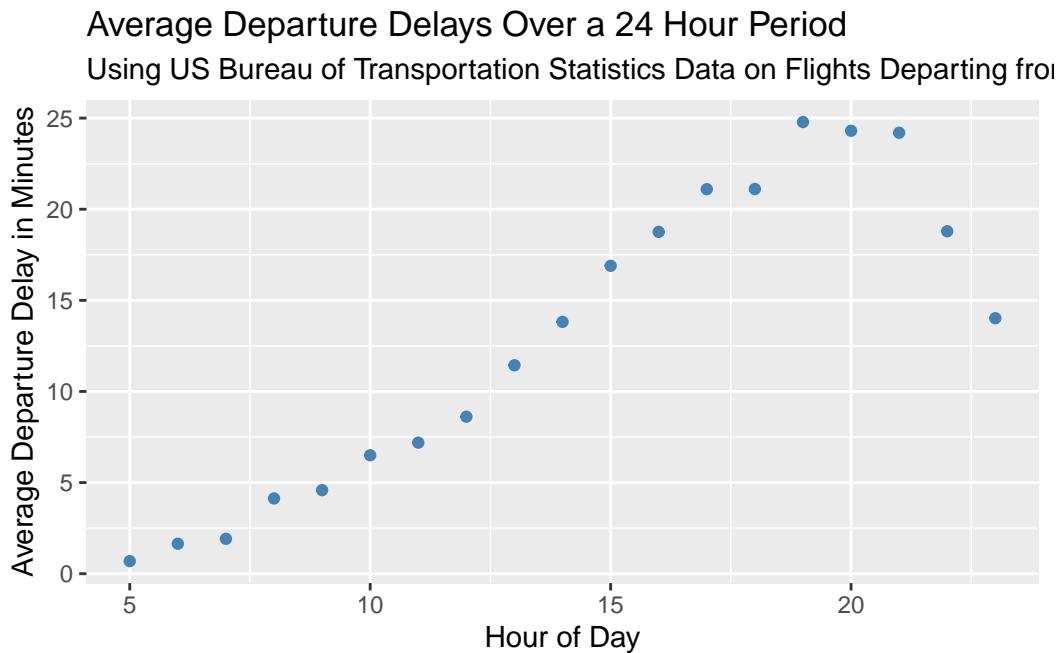
```

The carrier with the worst average delays is Frontier airlines. However, they are headquartered in Denver and this data set only includes flights from LGA to DEN. This means there's no variance in F9 across airports since there is only one flight path being recorded. We cannot disentangle the effects of bad airports from this data set and would have to make a lot of assumptions in order to do so.

2. Find the flights that are most delayed upon departure to each destination.

```
# A tibble: 105 x 19
# Groups:   dest [105]
  dest carrier dep_delay year month day dep_time sched_dep_time arr_time
  <chr> <chr>     <dbl> <int> <int> <int>    <int>          <int>    <int>
1 HNL  HA        1301  2013     1     9      641         900    1242
2 CMH  MQ        1137  2013     6    15     1432        1935    1607
3 ORD  MQ        1126  2013     1    10     1121        1635    1239
4 SFO  AA        1014  2013     9    20     1139        1845    1457
5 CVG  MQ        1005  2013     7    22      845        1600    1044
6 TPA  DL        960   2013     4    10     1100        1900    1342
7 MSP  DL        911   2013     3    17     2321        810     135
8 PDX  DL        899   2013     6    27      959        1900    1236
9 ATL  DL        898   2013     7    22     2257        759     121
10 MIA  AA       896   2013    12     5      756        1700    1058
# i 95 more rows
# i 10 more variables: sched_arr_time <int>, arr_delay <dbl>, flight <int>,
# tailnum <chr>, origin <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
# minute <dbl>, time_hour <dttm>
```

3. How do delays vary over the course of the day? Illustrate your answer with a plot.



4. What happens if you supply a negative n to slice\_min() and friends?

```

# A tibble: 336,776 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>    <int>          <int>     <dbl>    <int>          <int>
1 2013     12      7    2040            2123     -43        40          2352
2 2013      2      3    2022            2055     -33       2240          2338
3 2013     11     10    1408            1440     -32       1549         1559
4 2013      1     11    1900            1930     -30       2233         2243
5 2013      1     29    1703            1730     -27       1947         1957
6 2013      8      9     729             755     -26       1002         955
7 2013     10     23    1907            1932     -25       2143         2143
8 2013      3     30    2030            2055     -25       2213         2250
9 2013      3      2    1431            1455     -24       1601         1631
10 2013      5      5     934             958     -24       1225         1309
# i 336,766 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dttm>

# A tibble: 336,776 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>    <int>          <int>     <dbl>    <int>          <int>
1 2013     12      7    2040            2123     -43        40          2352
2 2013      2      3    2022            2055     -33       2240          2338
3 2013     11     10    1408            1440     -32       1549         1559
4 2013      1     11    1900            1930     -30       2233         2243
5 2013      1     29    1703            1730     -27       1947         1957
6 2013      8      9     729             755     -26       1002         955
7 2013     10     23    1907            1932     -25       2143         2143
8 2013      3     30    2030            2055     -25       2213         2250
9 2013      3      2    1431            1455     -24       1601         1631
10 2013      5      5     934             958     -24       1225         1309
# i 336,766 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dttm>

```

The result is listing the rows based on the smallest to largest dep\_delay value. This result is the same as arranging the values by dep\_delay.

5. Explain what count() does in terms of the dplyr verbs you just learned. What does the sort argument to count() do?

```
# A tibble: 3 x 2
  origin      n
  <chr>    <int>
1 EWR      120835
2 JFK      111279
3 LGA      104662
```

```
# A tibble: 3 x 2
  origin      n
  <chr>    <int>
1 EWR      120835
2 JFK      111279
3 LGA      104662
```

Count allows you to count the unique values of one or more variables similar to group\_by and summarise. The sort options causes the results to be presented in descending order.

6. Suppose we have the following tiny data frame:

- Write down what you think the output will look like, then check if you were correct, and describe what group\_by() does.

```
# A tibble: 5 x 3
# Groups:   y [2]
  x y     z
  <int> <chr> <chr>
1 1 a     K
2 2 b     K
3 3 a     L
4 4 a     L
5 5 b     K
```

This command should group by the two groups in y.

- Write down what you think the output will look like, then check if you were correct, and describe what arrange() does. Also, comment on how it's different from the group\_by() in part (a).

```
# A tibble: 5 x 3
  x y     z
  <int> <chr> <chr>
1 1 a     K
```

```

2      3 a      L
3      4 a      L
4      2 b      K
5      5 b      K

```

This command should arrange the results by the y column alphabetically. This means it should be list the results as a,a,a,b,b. The row order changed in using `arrange()` and not `group_by()`.

- c. Write down what you think the output will look like, then check if you were correct, and describe what the pipeline does.

```

# A tibble: 2 x 2
  y      mean_x
  <chr>   <dbl>
1 a          2.67
2 b          3.5

```

The result should be a summary of the mean of x based on the two groups of y. The pipeline passes the `group_by()` function along into the `summarize()` function so the means are groups by y values.

- d. Write down what you think the output will look like, then check if you were correct, and describe what the pipeline does. Then, comment on what the message says.

```
`summarise()` has grouped output by 'y'. You can override using the `.`groups` argument.
```

```

# A tibble: 3 x 3
# Groups:   y [2]
  y      z      mean_x
  <chr> <chr>   <dbl>
1 a      K          1
2 a      L          3.5
3 b      K          3.5

```

The result should be a summary of the mean of x based on the three groups of y and z (a K, a L, a K). This pipeline passes the `group_by()` function along into the `summarize()` function so the means are groups by y and z values. The message means that the results are grouped by y values first and the grouping can be overriden using the `.groups` argument.

- e. Write down what you think the output will look like, then check if you were correct, and describe what the pipeline does. How is the output different from the one in part (d)?

```
# A tibble: 3 x 3
  y      z    mean_x
  <chr> <chr>   <dbl>
1 a      K        1
2 a      L        3.5
3 b      K        3.5
```

The answer should be the same except the results are not grouped.

- f. Write down what you think the outputs will look like, then check if you were correct, and describe what each pipeline does. How are the outputs of the two pipelines different?

``summarise()` has grouped output by 'y'. You can override using the `~.groups` argument.`

```
# A tibble: 3 x 3
# Groups:   y [2]
  y      z    mean_x
  <chr> <chr>   <dbl>
1 a      K        1
2 a      L        3.5
3 b      K        3.5

# A tibble: 5 x 4
# Groups:   y, z [3]
  x     y      z    mean_x
  <int> <chr> <chr>   <dbl>
1 1     a      K        1
2 2     b      K        3.5
3 3     a      L        3.5
4 4     a      L        3.5
5 5     b      K        3.5
```

The first function calculated the means for the three combinations of the y and z column values. The second function creates a new column where these means are stored. The output of the first function is three rows and five rows for the second function.